



NUI Galway  
OÉ Gaillimh

# Machine Learning

## Week 10: Linear Classifiers with Hard and Soft Thresholds

Prof. Michael Madden

Chair of Computer Science

Head of Machine Learning & Data Mining Group

National University of Ireland Galway



# Learning Objectives

After successfully completing this topic, you will be able to ...

- Explain the drawbacks of using linear regression directly for classification problems
- Describe and implement approaches to improve on this:  
Linear Classifiers and Logistic Regression
- Discuss their characteristics and limitations



# Classification

- Is unknown sample  $\times$  or  $\circ$  ?

Goal: Find model that correctly classifies a new sample,  $\mathbf{x}_i$

- Training Data:  $(\mathbf{x}_i : y_i)$

$s_1: (11, 3 : +1) \circ$

$s_2: (3, 10 : +1) \circ$


$s_3: (14, 17 : -1) \times$

$s_4: (16, 16 : -1) \times$

...

- Unknown Sample  $s = (3, 15 : ?)$


What class does this belong to?







# From Linear Regression to Classification (1)

- Suppose we want to build a Yes/No **classification** model
- We know how to do linear regression:
  - Could encode No as 0 and Yes as 1
  - Perform linear regression to find  $h_{\theta}(x)$
  - Threshold predictions:  
 $h_{\theta}(x) \geq 0.5 \Rightarrow \text{Yes}$   
 $h_{\theta}(x) < 0.5 \Rightarrow \text{No}$
- Unfortunately, using Linear Regression directly doesn't always work very well...





## From Linear Regression to Classification (2)







# From Linear Regression to Classification (3)

- The reason for the problem:

Linear regression parameters are found without taking into account that there are only two 'real' output values, 0/1

A threshold is applied to outputs to convert them to 0/1, but **only after** the regression hyperplane is learned

- The solution:

Incorporate the threshold in the objective function, so that it is taken into account in the cost function and therefore **becomes part of the objective to be learned**


Could use a **hard** or **soft** threshold:

lead to **Linear Classifiers** & **Logistic Regression**, respectively...




# Linear Classifier

Goal: Given sample  $\mathbf{x}$ ,  
find function  $f(\mathbf{x})$  that correctly classifies it



$$f(\mathbf{x}) = \theta \cdot \mathbf{x}$$

$f(x) < 0.5$ : sample is  $\mathbf{x}$   
 $f(x) \geq 0.5$ : sample is  $\mathbf{o}$



As before, add a dummy input variable  $x_0 = 1$

To classify a new sample,  $\mathbf{x}$ :


- Calculate the dot product of the weight vector,  $\theta$ , and  $\mathbf{x}$
- Apply a hard threshold:

$$Threshold(z) = 1 \text{ if } z \geq 0.5, 0 \text{ otherwise}$$



# Building a Linear Classifier

- To build a linear classifier:  
$$h_{\theta}(\mathbf{x}) = \text{Threshold}(f(\mathbf{x})) = \text{Threshold}(\boldsymbol{\theta} \cdot \mathbf{x})$$
  - Find a function  $f(\mathbf{x})$  (i.e. find values for  $\boldsymbol{\theta}$ ) that correctly classifies training data **when put through threshold function**
- Can use Gradient Descent to find values for  $\boldsymbol{\theta}$ 
  - But  $h_{\theta}(\mathbf{x})$  is not differentiable, because of hard threshold function ...







# Building a Linear Classifier: Perceptron Learning Rule

- Because  $h_{\theta}(\mathbf{x})$  is not differentiable, cannot use the exact same approach that we used for Linear Regression
  - Instead need **Perceptron Learning Rule** to update  $\theta$  values
    - Also called other names
    - Usually used with Stochastic Gradient Descent
- $$\theta_j \leftarrow \theta_j + \alpha(y - h_{\theta}(\mathbf{x})) \cdot x_j \quad \text{for a single training case } (\mathbf{x}, y)$$
- Notes:
    - For Stochastic GD, picking only one sample, so  $N=1$
    - Converges to a solution, provided data linearly separable



# How does Perceptron Learning Rule Compare to Linear Regression Update Rule?

- **Perceptron Learning Rule:**

$$\theta_j \leftarrow \theta_j + \alpha(y - h_{\theta}(\mathbf{x})) \cdot x_j \quad \text{for a single example } (\mathbf{x}, y)$$

- Linear Regression update rule from last topic:

$$\theta_j \leftarrow \theta_j - \alpha(h_{\theta}(\mathbf{x}) - y)x_j$$

Set  $N=1$ , merge equations ...



# Building a Linear Classifier: Perceptron Learning Rule

- Looks just like MLR update rule (previous topic), but behaviour is different as  $h()$  and  $y$  are 0 or 1:
  - Output correct ( $h_{\theta}(\mathbf{x}) = y$ ): weights unchanged
  - $y = 1$  but  $h_{\theta}(\mathbf{x}) = 0$ : increase  $\theta_i$  if  $x_i$  positive
  - $y = 0$  but  $h_{\theta}(\mathbf{x}) = 1$ : decrease  $\theta_i$  if  $x_i$  positive
- To guarantee convergence, need to **decay**  $\alpha$  in proportion to  $1/t$ 
  - For Perceptron, must use smaller values of  $\alpha$  in each successive iteration
  - Unlike Linear Regression update rule: fixed  $\alpha$  because gradients get smaller
  - $t$  is iteration (“time step”)

$$\theta_j \leftarrow \theta_j + \alpha(y - h_{\theta}(\mathbf{x})) \cdot x_j$$



# Putting it Together: Linear Classifier Using Stochastic GD with Perceptron Learning Rule

Linear Classifier Learning Algorithm:

**initialise**  $\theta$  to any set of valid initial values

**Initialise**  $\alpha_0$  to some step size

**repeat**  $t=1:T$  times, or until convergence if earlier:

**select** a training example at random

$$\alpha \leftarrow \alpha_0 / t$$

**simultaneously foreach**  $\theta_j$  in  $\theta$  **do**:


$$\theta_j \leftarrow \theta_j + \alpha(y - h_{\theta}(\mathbf{x})) \cdot x_j$$



# Linear Classifier: Result

- Example of a successful classifier
- Data are linearly separable: can perfectly classify them
- To classify unknown sample  
 $x = (3, 15: ?)$ :  
 $f(x) = 3 + 15 - 16$   
 $f(x) = 2$   
 $f(x) > 0 \Rightarrow$  sample is **X**
- Another unknown sample  
Green star:  $x = (15, 0)$   
 $f(x) = -1 \Rightarrow$  sample is **O**
- However, many functions could separate this data....


★ = Unknown Sample





# Linear Classifier: Which is Best?

- There are many linear classifiers to choose from
  - Which one you find will depend on parameter search settings
- All work equally on linearly separable training data
  - But some will output a different prediction for unknown samples
  - Because Perceptron Rule works with 1/0 values, converges fully as soon boundary line found to fully separate data
  - Where it stops depends on if approaching from “red side” or “blue side”







## Avoiding This Behaviour ...

- Would like to find a boundary line that falls between the two classes, separating them as well as possible
  - To address this, need a different approach:  
**Linear Support Vector Machine**
  - Finds the maximum margin hyperplane separating classes
  - Not within the scope of this module
- The Logistic Classifier (coming up next) uses a “soft margin” that tends to push boundary away from nearest training data
  - However Logistic Classifier is **not** guaranteed to maximize the separating plane, whereas SVM is guaranteed to.



# Soft Thresholds & Logistic Regression

- Instead of hard threshold used in Linear Classifier, could use a **soft** one

Allow  $h_{\theta}(x)$  to take on values in range  $[0,1]$

Have it switch rapidly from 0 to 1 (almost step function)


- Go from the linear regression formula:

$$h_{\theta}(x) = \theta \cdot x$$

- To this:

$$h_{\theta}(x) = g(\theta \cdot x) \quad \text{where} \quad g(z) = \frac{1}{1 + e^{-z}}$$

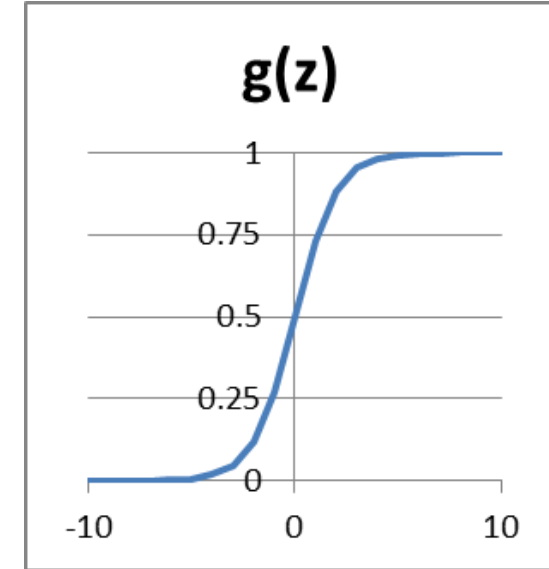
- $g(z)$  is called the sigmoidal or logistic function





# Logistic Regression

- How we interpret the output:  
 $h_{\theta}(\mathbf{x})$  is an estimate of the **probability** that  $y=1$  for input  $\mathbf{x}$ , given the parameters  $\theta$
- As before, can derive a cost function for  $h_{\theta}(\mathbf{x})$  and optimise parameters with gradient descent
  - The cost function is differentiable: can use standard GD as with Linear Regression
- Important: Logistic Regression is used for classification tasks, **not** for regression tasks!





# Logistic Regression Cost Function [1]

- Probability that  $y=1$  (Positive Class) for a case  $x$  is given by  $h_{\theta}(x)$

$$P(y=1 \mid x) = h_{\theta}(x)$$

- Therefore, probability that  $y=0$  (Negative class) is  $1 - h_{\theta}(x)$

$$P(y=0 \mid x) = 1 - h_{\theta}(x)$$

- We can combine these equations to cover both  $y=1$  and  $y=0$ :

$$P(y \mid x) = (h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$$

- Starting from this, a cost function can be defined, though I won't show its derivation (I include the index  $(i)$  for a training instance):

$$J(\theta) = -\frac{1}{N} \sum_{i=1}^N y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))$$





# Logistic Regression Cost Function [2]

$$J(\boldsymbol{\theta}) = -\frac{1}{N} \sum_{i=1}^N y^{(i)} \log(h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}))$$

- Behaviour:
  - As  $h_{\boldsymbol{\theta}}(\mathbf{x})$  tends to the correct value (either  $y=1$  or  $y=0$ ),  $J(\boldsymbol{\theta})$  tends to 0
  - As  $h_{\boldsymbol{\theta}}(\mathbf{x})$  tends to the wrong value,  $J(\boldsymbol{\theta})$  tends towards infinity
- The partial derivative of this cost function is:

$$\frac{\partial}{\partial \theta_j} J(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^N (h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}$$

- Surprisingly, this looks identical to the linear regression case, though the hypothesis function is different
- Note: There are various definitions and derivations; I am following Stanford ones: <http://ufldl.stanford.edu/tutorial/supervised/LogisticRegression/>





# Comparing Linear Classifiers With Hard and Logistic Thresholds

- Both find a hyperplane between classes, with threshold function to convert real number to 0/1
  - Both assume that classes are linearly separable
  - Neither attempt to maximise margin, though sigmoid can push Logistic Regressor boundary out from cases closest to it
- Parameters found with Gradient Descent
  - Logistic: Standard GD approach
  - Hard: Use a different update rule (Perceptron Update Rule) and have to decay  $\alpha$
- Logistic Regression outputs probabilities
  - Reflects uncertainty close to decision boundaries
- Logistic Regression has better convergence and behaves better when data are not linearly separable





# Learning Objectives Review

In this topic you have learned to ...

- Explain the drawbacks of using linear regression directly for classification problems
- Describe and implement approaches to correct this:  
Linear Classifiers and Logistic Regression
- Discuss their characteristics and limitations

**Final Note – these are important foundational concepts:  
Limitations of Linear Classifiers addressed with SVMs;  
Logistic Regression leads to ideas in Neural Networks.**